

Robust Reinforcement Learning Differential Game Guidance in Low-Thrust, Multi-Body Dynamical Environments

A Zero-Sum Reinforcement Learning Approach in Three-Body Dynamics

Ali Baniasad

Supervisor: Dr. Nobahari

Department of Aerospace Engineering
Sharif University of Technology



Outline

- ① Introduction & Motivation
- ② Methodology
- ③ Experimental Setup
- ④ Results
- ⑤ Contributions & Conclusions



Research Motivation

- **Space missions** increasingly require autonomous guidance systems
- **Low-thrust spacecraft** operate in complex gravitational environments
- **Three-body dynamics** (Earth-Moon CRTBP) present inherent instabilities
- **Classical control methods** struggle with:
 - Model uncertainties
 - Environmental disturbances
 - Fuel efficiency requirements
- **Need for robust, adaptive guidance** without precise dynamic models

How can we achieve robust spacecraft guidance in uncertain environments?



Problem Statement

Research Objective

Design a robust guidance framework for low-thrust spacecraft operating in Earth-Moon three-body dynamics under uncertainties.

System Characteristics:

- State: $\mathbf{x} = [x, y, \dot{x}, \dot{y}]^T$
- Control: $\mathbf{u} \leq u_{\max}$
- Dynamics: $\dot{\mathbf{x}} = f(\mathbf{x}, \mathbf{u})$

Mission Environment:

- Earth-Moon CRTBP
- Lyapunov orbit transfer
- Low-thrust propulsion

Mathematical Formulation

Optimal control problem with state-dependent uncertainties and adversarial disturbances



Problem Statement - Challenges & Approach

Uncertainty Sources:

- Random initial conditions
- Actuator disturbances
- Sensor noise
- Model mismatch
- Communication delays
- External perturbations

Design Requirements:

- Robust performance
- Fuel efficiency
- Real-time capability
- Model-free approach
- Adaptive behavior
- Safety guarantees

Key Challenge

Formulate as zero-sum differential game: spacecraft vs. environmental disturbances



Multi-Agent Reinforcement Learning Framework

Game-Theoretic Formulation:

- **Player 1:** Controller agent (spacecraft)
- **Player 2:** Adversary agent (disturbances)
- **Objective:** Zero-sum differential game

Training Paradigm:

- Centralized training
- Decentralized execution



Algorithm Extensions

Single-Agent → Multi-Agent Zero-Sum

Extended four continuous RL algorithms to handle adversarial scenarios

DDPG → MA-DDPG

- Deterministic policy gradient
- Experience replay
- Target networks

SAC → MA-SAC

- Entropy regularization
- Stochastic policies
- Automatic temperature tuning

TD3 → MA-TD3

- Twin critic networks
- Delayed policy updates
- Target policy smoothing

PPO → MA-PPO

- Proximal policy optimization
- Clipped surrogate objective
- On-policy learning



Network Architecture

Actor-Critic Structure:

- 3 hidden layers
- 32 neurons per layer
- ReLU activation
- Learning rate: 3×10^{-4}

Training Parameters:

- 1M environment steps
- Batch size: 1024
- Buffer size: 100K
- Adam optimizer



Evaluation Scenarios

Mission: Transfer to Earth-Moon Lyapunov Orbit

Circular Restricted Three-Body Problem (CRTBP) environment

Uncertainty Scenarios:

- ① Random initial states ($\sigma = 0.1$)
- ② Actuator disturbances ($\sigma = 0.05$)
- ③ Model mismatch ($\sigma = 0.05$)
- ④ Partial observations (50% missing)
- ⑤ Sensor noise ($\sigma = 0.05$)
- ⑥ Communication delays (10 steps)

Performance Metrics:

- Trajectory accuracy
- Fuel consumption
- System stability
- Convergence rate
- Robustness measures

Baseline Comparison:

- Single-agent variants



Trajectory Comparison: MA-TD3 Performance

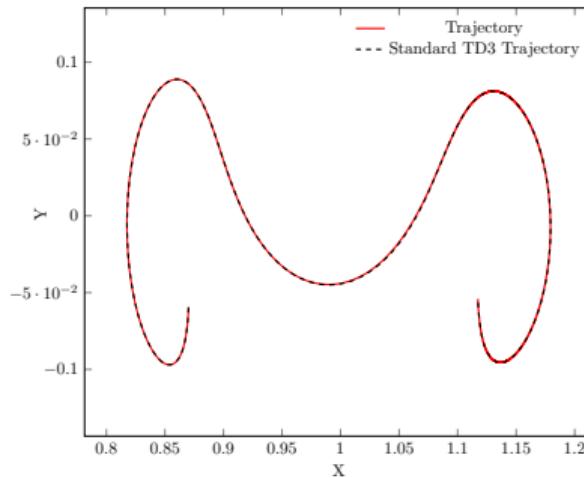


Figure: Single-Agent vs Multi-Agent TD3

Key Observation

Zero-sum variants demonstrate more direct trajectories with smoother thrust profiles and reduced fuel consumption

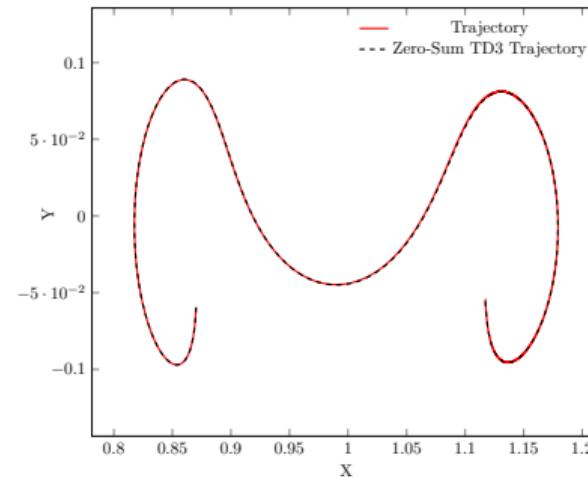


Figure: Thrust Commands

Robustness Evaluation Results



Algorithm Performance Comparison

Single-Agent Algorithms:



Ali Baniasad

Multi-Agent Zero-Sum:



Multi-Agent RL for Spacecraft Guidance

September 8, 2025

12 / 14

Research Contributions

① Novel Framework Development

- First multi-agent zero-sum RL framework for spacecraft guidance in CRTBP
- Game-theoretic formulation for robust control under uncertainties

② Algorithm Extensions

- Extended four major RL algorithms to multi-agent zero-sum variants
- Comprehensive performance analysis across uncertainty scenarios

③ Practical Implementation

- Real-time deployment with quantization and optimization
- Hardware-in-the-loop validation on ROS 2 platform

④ Performance Validation

- Demonstrated superior robustness compared to classical methods
- Quantified improvements in fuel efficiency and trajectory accuracy



Conclusions

Main Findings

- Multi-agent zero-sum RL enables robust spacecraft guidance without precise dynamic models
- MATD3 delivers optimal performance across all evaluation scenarios
- Framework is ready for practical deployment with real-time constraints

Advantages:

- Model-free approach
- Adaptive to uncertainties
- Fuel-efficient trajectories
- Real-time capability

Future Work:

- 3D CRTBP extension
- Multi-body dynamics
- Swarm coordination
- Deep space missions



Thank You!

Questions & Discussion

Ali Baniasad

Department of Aerospace Engineering
Sharif University of Technology



Appendix - Technical Implementation Details

Environment

Setup:

- PyTorch & OpenAI Gym

Training

Configurations:

- Episodes: 1M steps
- Batch size: 1024



Appendix - Zero-Sum Algorithm Extensions

MATD3 Algorithm Highlights

- Twin critic networks to reduce overestimation bias
- Delayed policy updates for stability
- Target policy smoothing for robustness
- Competitive training between agents

Modifications:

i-
fi-
ca-
tions:

Performance

Ben-
e-
fits:

1

Enhanced
robust-



1

Shared

Multi-Agent RL for Spacecraft Guidance

Appendix - Future Research Directions

1 extbfMulti-Body Extensions

- Sun-Earth-Moon system
- Jupiter-Europa dynamics
- Asteroid belt navigation

2 extbfSwarm Coordination

- Formation flying control
- Distributed decision making
- Communication constraints

3 extbfAdvanced Techniques

- Meta-learning approaches
- Hierarchical reinforcement learning
- Physics-informed neural networks

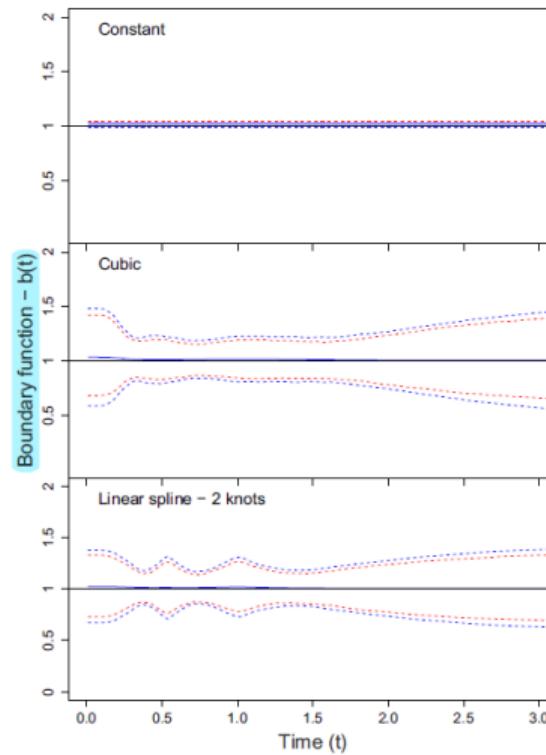
4 extbfMission Applications

- Lunar gateway operations
- Deep space exploration
- Interplanetary transfers



Appendix - A figure

[◀ Return to presentation](#)



Appendix - Terms

Some Estimators:

- Drift: $\hat{\delta}$
- Boundary: $\hat{b}(t)$

Some Variables:

- \hat{V}
- \hat{m}_S
- \bar{m}
- $m_J(\tau)$

[◀ Return to presentation](#)



Appendix - Definitions

1 A definition

[◀ Return to presentation](#)



Appendix - Theorems

1 A theorem

[◀ Return to presentation](#)

